Accelerating app modernisation with New Relic AI Monitoring

Phil Weber
New Relic University

© 2025 New Relic, Inc. All rights reserved.

Agenda

01	Introduction
02	New Relic AI vs AI Monitoring
03	How it works
04	Demo and Lab

© 2025 New Relic, Inc. All rights reserved.

Do you ask these questions?

- Do I have AI monitoring gaps?
- What AI requests are being sent?
- How is the performance of my AI response(s)?
- Are there any issues or errors?
- Is performance affecting users / user experience?

© 2025 New Relic, Inc. All rights reserved.

new relic

Challenges monitoring AI applications

Performance Tuning

AI tech stacks include new components like LLMs, vector data stores, and orchestration frameworks

Quality Control

AI applications introduce new challenges with quality of responses (bias, hallucination and toxicity)

Manage Cost

Managing costs associated with model usage and resource utilization is a challenge

Security

Customers don't trust sharing proprietary information with AI/LLM providers

Compliance

AI regulations are still emerging and customers need to audit and prove compliance

© 2025 New Relic, Inc. All rights reserved.

new relic

New Relic AI Monitoring (AIM)

- **Quick and easy setup**
  New Relic agents come fully loaded with all the AIM capabilities
- **Debug faster with complete visibility**
  End-to-end visibility of your entire AI application stack
- **Ensure AI app performance, quality, and cost**
  Follow the end-to-end lifecycle of LLM prompts and responses to optimize for latency, quality, and cost of AI applications
- **Choose the right model for your app**
  Track usage, performance, quality, and cost across all models in a single view to choose the optimal model for your needs
- **Instantly monitor any AI ecosystem**
  50+ integrations for popular LLMs, vector data stores, and more

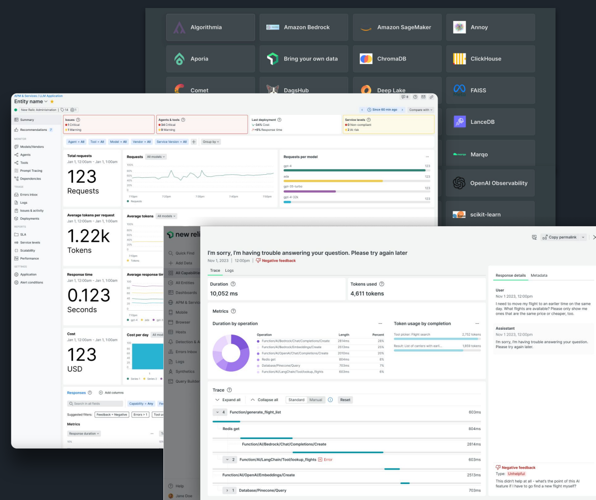

© 2025 New Relic, Inc. All rights reserved.

new relic

New Relic AI vs AIM

New Relic AI

New Relic AI combines large language models, specialized tools, and our consolidated platform to make observability easier for any user, from observability novices to seasoned practitioners.

New Relic AI Monitoring

Industry's first end-to-end APM solution that brings the power of AI observability to help engineers build and run AI applications confidently with in-depth insights to optimize performance, quality and cost.

© 2025 New Relic, Inc. All rights reserved.

new relic

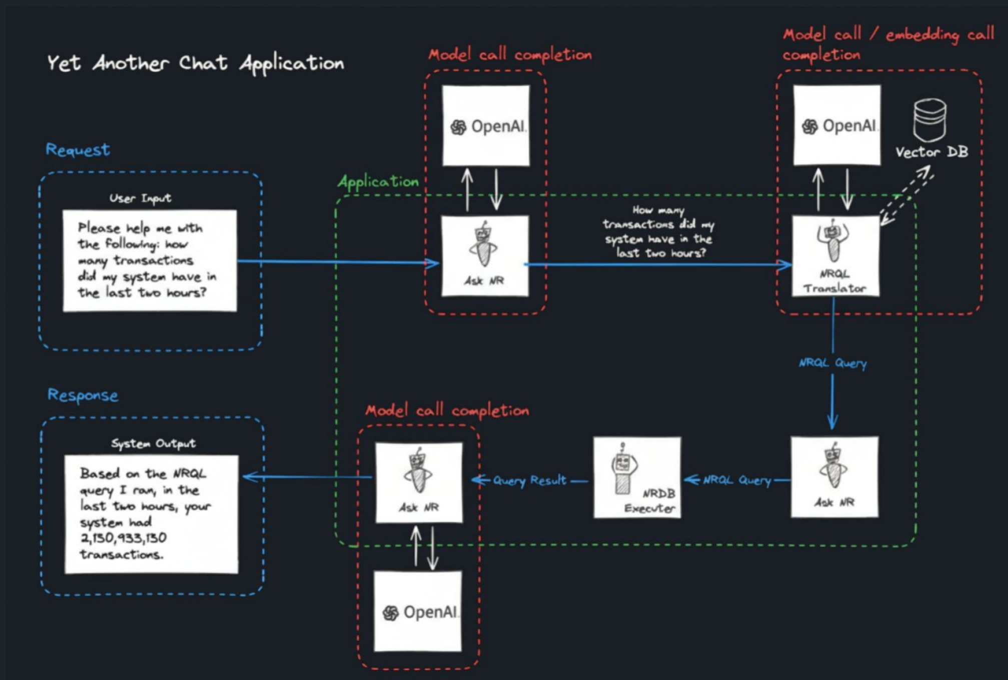

new relic

How it works

- Built into APM agents, just enable it
- Monitors the AI black box used by your app
- Records model performance, token usage, and completions
- Gathers metrics and events from LLM requests and responses

© 2025 New Relic, Inc. All rights reserved.

new relic

New Relic APM with AI Monitoring

Performance

- Are users waiting too long for a response?
- Spike in token responses?
- Negative user feedback around certain topics?
- AI responses included in trace data

Model Comparison

- Records performance and completion data of all AI models
- Isolate token usage
- Analyze performance of multiple models over time

© 2025 New Relic, Inc. All rights reserved.

new relic

New NRDB events

LlmChatCompletionMessage

Corresponds to each message (sent and received) from a chat completion call. Includes messages created by the user, assistant, and system.

LlmChatCompletionSummary

Captures high-level data about the creation of a chat completion for request and response messages.

LlmEmbedding

Captures data specific to the creation of an embedding.

https://docs.newrelic.com/attribute-dictionary/?dataSource=APM

© 2025 New Relic, Inc. All rights reserved. Confidential and proprietary. For internal use only, not for external distribution.

new relic

Lab: Using New Relic AIM

Select the **chat-service-node** service in the **Contributor Apps** account and use New Relic AIM to answer the following questions:

On the *AI Responses* page:

- What is the 95th percentile response time for the past 24 hours?
- Considering only requests LIKE %chess%, what is the average token usage per request?
- What is the total number of error responses in the past 3 days?

© 2025 New Relic, Inc. All rights reserved.

new relic

Lab: Using New Relic AIM

Select the **chat-service-node** service in the **Contributor Apps** account and use New Relic AIM to answer the following questions:

On the *Model Inventory* page:

- What are the four tabs available for model analysis?
- How is *Cost* measured?
- What metrics are shown on the *Performance* page?

© 2025 New Relic, Inc. All rights reserved.

new relic